

Reliability-weighted target prioritization in CD4+ T-cell Perturb-seq: a generalizability-theory decomposition

Che Cheng¹

¹Institute of Statistical Science, Academia Sinica, Taipei, Taiwan.
che830621@stat.sinica.edu.tw. ORCID 0000-0003-3376-7833.

July 12, 2026

Abstract

Genome-wide Perturb-seq screens prioritize candidate targets by the strength of a perturbation’s transcriptional effect. Effect strength does not answer a prior measurement question: is the readout dependable? A large effect estimated from a single guide, a single donor, or a pseudobulk of few cells need not survive replication, and for target prioritization each false lead costs a validation experiment. We treat each perturbation effect as a measurement in a crossed Target $\times$ Guide $\times$ Donor $\times$ Condition design and apply generalizability theory (Cronbach et al., 1972) to separate the dependable part of an effect from facet-specific idiosyncrasy. Guides and donors enter as random facets; condition enters as a fixed facet and is analyzed within its levels. For each target we report a dependability profile over the facets and a joint generalizability coefficient over the two random facets, and we re-rank targets by effect magnitude weighted by that coefficient. On the released screen (Zhu et al., 2025), removing the measurement-error floor estimated from the non-targeting controls raises the number of genes with a dependable target-signal share above .10 from 40 to 7,674. A design study indicates that reliability is limited by the number of guides rather than the number of donors. Analyzed within activation states, dependability recovers the T-cell-receptor signaling module as reliably measurable only in activated cells, without recourse to gene annotation. Every methodological decision was recorded and adversarially reviewed, and all results regenerate from the released summary statistics.

1 Dependability, not effect size

Genome-scale Perturb-seq couples pooled CRISPR perturbation with a single-cell transcriptomic readout, so that the transcriptional consequence of perturbing each of thousands of genes is measured in one experiment (Dixit et al., 2016; Replogle et al., 2022). Applied to primary human CD4+ T cells, the cells that coordinate adaptive immunity and are a substrate for engineered cell therapies, such a screen identifies the genes whose perturbation most reshapes a T-cell program and, on that basis, nominates candidate targets (Zhu et al., 2025). The nomination carries a cost. Each target that advances is committed to an arrayed knockdown, a functional assay, or an animal model, so a hit that does not reproduce spends a validation experiment on an artifact. What a screen is worth therefore depends not only on the effects it detects but on whether the ranking it induces over targets can be trusted.

Target prioritization in Perturb-seq conventionally ranks candidates by differential-expression strength: effect size, z -score, false-discovery rate, or the number of downstream differentially expressed genes. These summaries answer whether a perturbation moved the transcriptome *in the observed dataset*. They do not answer whether the effect would hold up under another guide, another donor, or another cell-state context. The two questions come apart. A target can carry a large apparent effect that is fragile, when two guides disagree, one donor drives the signal, knockdown is weak, or an off-target flag is present; a target of moderate effect that is consistent across guides and donors can be the better validation bet. We therefore treat *dependability* as a first-class quantity alongside magnitude.

Two properties are at stake that effect size does not by itself deliver: *accuracy*, that a reported effect is dependable rather than measurement noise, and *generalizability*, that it recurs across guides, donors, and cell-state contexts. We therefore ask three questions, at two levels of the measurement. *RQ1 (per target)*: which targets carry an individually dependable effect, one that would recur under a fresh guide and a fresh donor, rather than a large but fragile one? *RQ2 (the design)*: how dependable is the screen as a whole, at its two guides and four donors? *RQ3 (improving the design)*: which facet limits that dependability, so that a future screen knows whether to add guides or donors? RQ1 is the per-target coefficient on which we rank targets; RQ2 and RQ3 are design-level questions, answered by projecting the same variance decomposition to alternative designs. Answering any of them first requires separating the dependable part of an effect from measurement noise.

This concern is not new in high-throughput biology, though it has rarely been framed as measurement reliability. Li et al. (2011) call discoveries by their reproducibility across replicates rather than by effect size, defining an irreproducible discovery rate as an analogue of the false-discovery rate; Billmann et al. (2023) show that standard CRISPR-screen quality metrics do not by themselves capture the reproducibility of context-specific hits; and data-driven target-prioritization frameworks integrate several weighted evidence types rather than fitness effect alone (Behan et al., 2019). Generalizability theory supplies the missing piece: a single interpretable coefficient for the dependability of a per-target effect, together with an explicit account of which experimental facet limits it.

The rest of the paper develops the per-target coefficient and the two levels it rests on (Section 2), describes the data and the pre-registered choice of estimator (Section 3), and takes up RQ1 through RQ3 across four analyses (Section 4).

2 A generalizability-theory decomposition

2.1 The crossed design and its facets

We treat one perturbation’s effect as a measurement and ask how much of it would recur if the experiment were repeated. The quantity of interest is the target gene’s true effect on the transcriptome. What we observe is that effect seen through one guide, in one donor, under one culture condition, and read from a finite sample of cells, and each of these can make the observation differ from the truth. This is the setting generalizability theory was built for, and the correspondence is direct (Lord & Novick, 1968): the target gene is the *object of measurement*, its true effect is the *universe score* (the value one would obtain by averaging over infinitely many independent repeats), the two guides are interchangeable versions of the same perturbation (*parallel forms*), and donors and conditions are further experimental factors, or *facets*, crossed with the target. Cell sampling and sequencing contribute the residual noise. Classical test theory writes an observation as a true score plus error, $X = T + E$ (Novick, 1966), and generalizability theory (Cronbach et al., 1972)

splits that error into one term per facet,

$$E = E_{\text{guide}} + E_{\text{donor}} + E_{\text{condition}} + E_{\text{interactions}} + E_{\text{residual}},$$

and estimates how large each term is.

We distinguish random from fixed facets. Guides and donors are *random* facets: one wants the effect to generalize across the guides and donors that were not sampled. Condition is a *fixed* facet: Rest, Stim8hr, and Stim48hr are three deliberately chosen activation states, not a sample from a population of conditions, so the appropriate treatment is a separate analysis within each state rather than a generalization across states. A perturbation that acts only in the activated state is context-specific, not unreliable, and folding condition into a generalization coefficient would misclassify that signal as noise.

2.2 Per-target dependability

For target t we report dependability as a *profile* of three coefficients, one per facet, kept separate and never multiplied together. Each coefficient answers a reproducibility question: with everything held fixed except one facet, how well does the target’s transcriptional signature repeat? An observation is a (target, guide, donor, condition) pseudobulk whose effect profile spans the 18,129 genes, measured relative to a matched non-targeting control. We average a group of observations into a mean profile, and each facet coefficient is the Pearson correlation, across genes, between two mean profiles that differ only in that facet:

- R_t^{guide} : the target’s guides are split into two halves, each half’s mean profile is formed, and the two are correlated. A high value indicates that the transcriptional signature reproduces across independent guide sets.¹
- R_t^{donor} : the same construction with a two-versus-two donor split, then shrunk toward the pooled mean by a moment-based empirical-Bayes step, which stabilizes a coefficient that four donors otherwise leave with only three degrees of freedom.
- $R_t^{\text{condition}}$: the mean of the pairwise correlations among the three within-state profiles, reported as cross-context *consistency* rather than as a generalization coefficient.

To rank targets we need one number rather than a profile. The per-target dependability $R_{\text{dep},t}$ is that number, and it has a concrete meaning: the correlation one would expect between two fully independent repeats of the experiment, a second run carried out with fresh guides and fresh donors. Writing an observed profile as a true signal plus independent guide and donor errors, $\tau_t + \varepsilon_{\text{guide}} + \varepsilon_{\text{donor}}$, that correlation is the signal’s share of the total variation, $\sigma_\tau^2 / (\sigma_\tau^2 + \sigma_{\text{guide}}^2 + \sigma_{\text{donor}}^2)$. The same quantity is assembled from the two split-half coefficients. Writing each as $R = \tau / (\tau + e)$ for a shared signal variance τ and a facet error variance e , each facet contributes an error-to-signal ratio $(1 - R)/R = e/\tau$, and because the guide and donor errors are independent these ratios add:

$$R_{\text{dep},t} = \frac{1}{1 + \frac{1 - R_t^{\text{guide}}}{R_t^{\text{guide}}} + \frac{1 - R_t^{\text{donor}}}{R_t^{\text{donor}}}}, \quad S_t = E_t R_{\text{dep},t}, \quad (1)$$

¹These are raw single-observation split-half correlations. The reliability of the deployed two-guide average is the Spearman-Brown step-up $2r/(1+r)$, a separate and larger quantity; we report the single-observation form because the joint coefficient is defined at the single-guide-by-single-donor level. Both are distribution-light empirical correlations rather than Gaussian variance-component ratios.

where E_t is the effect magnitude and S_t the dependability-weighted score used for ranking. A single-facet coefficient refreshes one source of variation; the joint coefficient refreshes both, and is therefore the most demanding form of reproducibility we report. The equality is exact when the two errors are independent, and becomes an upper bound if a guide-by-donor interaction adds a further disagreement term.

We keep a naming discipline that separates the two levels of the theory: $R_{\text{dep},t}$ is a per-target quantity, whereas the design-level coefficients $E\rho^2$ and Φ describe the measurement design as a whole and carry no per-target subscript.

3 Data and estimation

The data are the released summary statistics of a genome-scale Perturb-seq screen in primary human CD4+ T cells (Zhu et al., 2025): CRISPR perturbations of 18,129 genes in a crossed design of two guides per gene, four donors, and three culture conditions (Rest, Stim8hr, Stim48hr), with non-targeting guides as controls. We work from the released joint pseudobulk and define each perturbation’s per-gene effect as its non-targeting-control-relative log-CPM within the matched donor-condition-run stratum, so that the effect is measured against the controls that share its technical context. No single-cell data are processed.

The variance contributed by each facet is estimated per gene by a crossed mixed model that weights each pseudobulk by its precision, following the generalizability-theory design of Section 2 (Brennan, 2001). We did not assume this estimator. Three candidates were compared on simulated data with known variance components, a pre-registered synthetic-recovery test, and only the precision-weighted model recovered the planted values under both Gaussian and heavy-tailed noise. A distance-based candidate failed, because a share of a distance is not a variance and cannot be read as one.

4 Results

4.1 Measurement error masks dependable signal

The residual of a pseudobulk decomposition combines biological idiosyncrasy with measurement error, and the latter is inflated when a pseudobulk aggregates few cells. The non-targeting controls carry no perturbation, so the variance of the non-targeting effect within its stratum estimates the per-gene measurement-error floor directly. Subtracting this floor from the residual raises the median target-signal share from .019 to .081 and raises the number of genes with a target-signal share above .10 from 40 to 7,674, or 42% of the genome (Figure 1). The dependable signal was present but masked; the same controls give a clean non-targeting-versus-non-targeting negative test ($r \approx 0$), so the correction introduces no signal where none exists.

4.2 A guide-limited design

A design study projects the generalizability coefficient onto alternative numbers of guides and donors (Figure 2). On the representative signal genes the current design (two guides, four donors) reaches $E\rho^2 \approx .44$. The guide error term exceeds the donor error term by a factor of about 2.4; adding one guide raises the coefficient about three times as much as adding one donor. With two guides the coefficient is capped near .53, and no number of donors reaches .70, whereas .70 is attainable with about 15 guides. Guide-to-guide variability, rather than donor heterogeneity, is the binding constraint, which is actionable guidance for the design of subsequent screens. Repeating the design

Removing measurement error via the non-targeting controls




Figure 1: Measurement-error removal via the non-targeting controls. Left, the median target-signal share; right, the number of genes with a dependable target-signal share above .10, on a logarithmic axis.

study within each activation state returns the same verdict in all three states (Table 1), so the recommendation is not an artifact of pooling; within a state the guide term dominates the donor term more sharply still, because most donor disagreement is a cross-state phenomenon.

4.3 A reliability-weighted ranking

Ranking targets by S_t reorders the candidate list relative to effect size alone (Figure 3). Of the 100 highest-effect targets, 49 fall out of the 100 most dependable, and reproducible targets that effect size underranks rise (Table 2). These are the targets a dependability-weighted screen would nominate for validation but a strength-ranked screen would overlook.

4.4 Context-specific dependability recovers the TCR module

Analyzed within activation states, dependability separates targets that are reliable across all states from targets that are reliable in one state and not another (Figure 4). Because a within-state coefficient rests on roughly a third of the cells and runs lower than the pooled value, states are compared to one another rather than to the pooled number. One hundred targets are context-specific in this sense. Among them, the T-cell-receptor (TCR) signaling module, the CD3 complex (CD3D, CD3G, CD247) with the proximal kinase ZAP70 and the scaffold LAT, is reliably measurable only in activated cells, with resting dependability at the noise floor (resting $R_{dep} \approx .02-.05$; stimulated $\approx .20-.28$). The knockdown effect for CD3D grows with activation (effect magnitude 0.26, 0.57, and 1.12

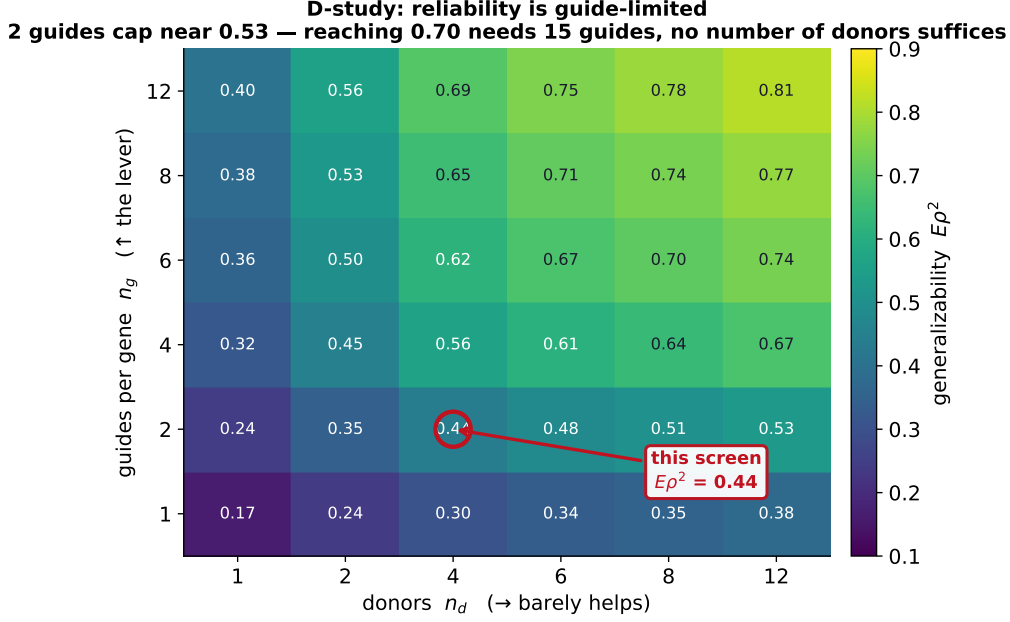


Figure 2: The generalizability coefficient $E\rho^2$ across the design grid. The current screen (two guides, four donors) is marked; reaching .70 requires about 15 guides and is unreachable by adding donors.

Table 1: Per-condition design study on the representative signal genes. The guide-limited verdict holds within every activation state.

Activation state	$E\rho^2$	Guide error	Donor error
Rest	.502	0.026	0.005
Stim8hr	.507	0.023	0.004
Stim48hr	.491	0.019	0.004
Pooled	.438	0.040	0.017

across Rest, Stim8hr, and Stim48hr) and becomes dependable only upon activation. The framework reconstructs a coherent signaling module from dependability alone, without gene annotation, which is the strongest available evidence that it measures a real property of the perturbation rather than an artifact of the summary.

5 Reproducibility

The pipeline runs end to end from the released summary statistics, without processing the full single-cell data: data retrieval, a checksummed evidence manifest, identifiability gates, synthetic method selection, per-gene decomposition, measurement-error removal, the design study, and the per-context ranking. All results are committed as tables and figures, and every reported number regenerates from a single re-run. The code is released under Apache-2.0.




Figure 3: Rank by effect size alone (horizontal) against rank by the dependability-weighted score (vertical). Points below the diagonal in the lower-left quadrant are demoted (strong but unreliable); points in the upper right are promoted (reproducible).

6 Limitations

Four donors leave the donor variance component with three degrees of freedom, so the donor coefficient and the donor axis of the design study carry real uncertainty, which we report as bands rather than point values. Sequencing run is partially confounded with donor, a fact of the released design that we record rather than correct. A small number of full-scale Monte-Carlo gate confirmations are scheduled on a compute cluster; the results shown are the local synthetic and real-data tier. Finally, this is a re-analysis of released data, so the design-study guidance is prospective, addressed to subsequent screens rather than to the present one.

A A recorded, cross-model-hardened process

The methodology was developed under an auditable process (Figure 5). Each decision (the choice of estimator, the thresholds, the classification of each facet as random or fixed, and the criterion for identifiability) was recorded as a versioned issue using issue-driven development (Cheng, 2026a) and spec-driven development, both open-source tooling the author maintains. Because the reasoning

Table 2: Targets that effect size underranks but dependability promotes. Rank is out of the 7,169 scored targets.

Target	Rank by effect	Rank by dependability	$R_{\text{dep},t}$
RPS3	79	4	.34
QARS1	53	5	.32
RPP21	124	8	.34
IMP4	96	13	.31
NMD3	158	17	.33

was documented rather than implicit, the frozen record could be subjected to adversarial review by an independent frontier model, which returned a blocked verdict with 11 findings, three of them critical, spanning identifiability, measurement error, selection, validation leakage, and compute. Each finding was resolved within the same recorded loop: the falsification gates and controls were frozen before any real result was inspected, the estimator was selected as in Section 3, and a mis-specified reliability aggregation (a product of coefficients) was replaced by the joint coefficient of Section 2.2. The resolutions were then cross-verified across models with a second adversarial pass and with `parallel-ai-agents` (Cheng, 2026b), a package the author maintains that dispatches a task to independent agents and cross-compares their outputs. The recorded decisions, the blocked verdict, and every resolution are versioned in the repository.

References

- Behan, F. M., et al. (2019). Prioritization of cancer therapeutic targets using CRISPR–Cas9 screens. *Nature*, 568, 511–516. <https://doi.org/10.1038/s41586-019-1103-9>
- Billmann, M., Ward, H. N., Aregger, M., Costanzo, M., Andrews, B. J., Boone, C., Moffat, J., & Myers, C. L. (2023). Reproducibility metrics for context-specific CRISPR screens. *Cell Systems*, 14(5), 418–422.e2. <https://doi.org/10.1016/j.cels.2023.04.003>
- Brennan, R. L. (2001). *Generalizability theory*. Springer.
- Cheng, C. (2026a). *Issue-driven development (IDD): A decision-ledger methodology for Claude Code*. <https://github.com/PsychQuant/issue-driven-development>
- Cheng, C. (2026b). *parallel-ai-agents: Parallel dispatch and cross-comparison across independent AI agents*. <https://github.com/PsychQuant/parallel-ai-agents>
- Cronbach, L. J., Gleser, G. C., Nanda, H., & Rajaratnam, N. (1972). *The dependability of behavioral measurements: Theory of generalizability for scores and profiles*. Wiley.
- Dixit, A., Parnas, O., Li, B., Chen, J., Fulco, C. P., Jerby-Arnon, L., Marjanovic, N. D., Dionne, D., Burks, T., Raychowdhury, R., Adamson, B., Norman, T. M., Lander, E. S., Weissman, J. S., Friedman, N., & Regev, A. (2016). Perturb-seq: Dissecting molecular circuits with scalable single-cell RNA profiling of pooled genetic screens. *Cell*, 167, 1853–1866.e17. <https://doi.org/10.1016/j.cell.2016.11.038>
- Li, Q., Brown, J. B., Huang, H., & Bickel, P. J. (2011). Measuring reproducibility of high-throughput experiments. *Annals of Applied Statistics*, 5(3), 1752–1779. <https://doi.org/10.1214/11-AOAS466>
- Lord, F. M., & Novick, M. R. (1968). *Statistical theories of mental test scores*. Addison-Wesley.
- Novick, M. R. (1966). The axioms and principal results of classical test theory. *Journal of Mathematical Psychology*, 3(1), 1–18. [https://doi.org/10.1016/0022-2496\(66\)90002-2](https://doi.org/10.1016/0022-2496(66)90002-2)





Figure 4: Per-state dependability for the T-cell-receptor signaling module. Resting values sit at the noise floor; the effect becomes reliably measurable only upon activation.

- Replogle, J. M., Saunders, R. A., Pogson, A. N., Hussmann, J. A., Lenail, A., Guna, A., Mascibroda, L., Wagner, E. J., Adelman, K., Lithwick-Yanai, G., Iremadze, N., Oberstrass, F., Lipson, D., Bonnar, J. L., Jost, M., Norman, T. M., & Weissman, J. S. (2022). Mapping information-rich genotype-phenotype landscapes with genome-scale perturb-seq. *Cell*, 185, 2559–2575.e28. <https://doi.org/10.1016/j.cell.2022.05.013>
- Zhu, R., Dann, E., Yan, J., Reyes Retana, J., Goto, R., Guitche, R. C., Petersen, L. K., Ota, M., Pritchard, J. K., & Marson, A. (2025). Genome-scale perturb-seq in primary human CD4+ T cells maps context-specific regulators of T cell programs and human immune traits. *bioRxiv*. <https://doi.org/10.64898/2025.12.23.696273>

An auditable, cross-model-hardened process



decide & document → get torn apart → repair (never “certified” — hardened)

Figure 5: The recorded, cross-model-hardened process: a decision ledger enables adversarial review by an independent model, whose blocked verdict is resolved in the same loop. The verdict was blocked and worked through, not certified.