

Reliability-weighted target prioritization in CD4+ T-cell Perturb-seq: a generalizability-theory decomposition

Che Cheng¹

¹Institute of Statistical Science, Academia Sinica, Taipei, Taiwan. che830621@as.edu.tw.
ORCID 0000-0003-3376-7833.

July 14, 2026

Abstract

Genome-scale Perturb-seq screens prioritize candidate targets by the strength of a perturbation’s transcriptional effect. Effect strength does not answer a prior measurement question: is the readout dependable? A large effect estimated from a single guide, a single donor, or a pseudobulk of few cells need not survive replication, and for target prioritization each false lead costs a validation experiment. We treat each perturbation effect as a measurement in a crossed Target $\times$ Guide $\times$ Donor $\times$ Condition design and apply generalizability theory (Cronbach et al., 1972) to separate the dependable part of an effect from facet-specific idiosyncrasy. Guides and donors enter as random facets; condition enters as a fixed facet and is analyzed within its levels. For each target we report a dependability profile over the facets and a joint generalizability coefficient over the two random facets, and we re-rank targets by effect magnitude weighted by that coefficient. On the released screen (Zhu et al., 2025), removing the measurement-error floor estimated from the non-targeting controls raises the number of genes with a dependable target-signal share above .10 from 40 to 7,674. Analyzed within activation states, dependability recovers the T-cell-receptor signaling module as reliably measurable only in activated cells, without recourse to gene annotation. A design study indicates that reliability is limited by the number of guides rather than the number of donors, so a future screen should add guides. Every methodological decision was recorded and adversarially reviewed, and all results regenerate from the released summary statistics.

1 Dependability in target prioritization

Genome-scale Perturb-seq couples pooled CRISPR perturbation with a single-cell transcriptomic readout, so that the transcriptional consequence of perturbing each of thousands of genes is measured in one experiment (Dixit et al., 2016; Replogle et al., 2022). Applied to primary human CD4+ T cells, the cells that coordinate adaptive immunity and are a substrate for engineered cell therapies, such a screen identifies the genes whose perturbation most reshapes a T-cell program and, on that basis, nominates candidate targets (Zhu et al., 2025). The nomination carries a cost. Each target that advances is committed to an arrayed knockdown, a functional assay, or an animal model, so a hit that does not reproduce spends a validation experiment on an artifact. What a screen is worth therefore depends not only on the effects it detects but on whether the ranking it induces over targets can be trusted.

Target prioritization conventionally ranks candidates by *effect size*, the strength of a perturbation’s differential-expression signature. The strength is read from the coefficients of a regularized model of expression on guide assignment (Dixit et al., 2016), from a *z*-score or false-discovery rate for each perturbation-gene association, or from the number of downstream differentially expressed genes together with a global test that the transcriptome moved at all (Replogle et al., 2022). Effect metrics have since grown more sophisticated, replacing the gene count with the distance between the perturbed and control single-cell distributions (Peidli et al., 2024) or with a per-target rate of cells classified as perturbed (Papalexi et al., 2021). All of them are measures of effect size, of how large an effect is.

None answers a prior question: whether the effect is dependable, in the sense that it would recur under another guide, another donor, or another cell-state context. The number of differentially expressed genes, for instance, tracks the number of cells assayed (Dixit et al., 2016) and need not track the strength of the perturbation itself (Norman et al., 2019), so a target can rank high on this measure yet low on reproducibility. A large apparent effect can be fragile, when two guides disagree, one donor drives the signal, knockdown is weak, or an off-target flag is present, while a target of moderate but consistent effect is the better validation bet. What decides between them is not which effect is larger in the observed data but which is more likely to reappear when the target is remeasured, because reappearing is what a validation experiment tests. Effect size and dependability can point to different targets, so a ranking that reads only the first keeps nominating the fragile ones. We therefore treat *dependability* as a first-class quantity alongside magnitude, and rank by the two together.

In generalizability theory (Cronbach et al., 1972), the psychometric account of when a measurement can be trusted, dependability has two senses. An effect is dependable when it is *accurate*, a real signal rather than measurement noise, and when it is *generalizable*, recurring across guides, donors, and cell-state contexts. The two show up the same way: an effect that is mostly noise, or that appears only with one guide or one donor, will not reproduce when it is measured again with fresh guides and donors. Measuring that *reproducibility* therefore captures both senses at once. We ask four questions, moving from the present screen to future ones:

- *RQ1 (per target)*: which targets carry a dependable effect, one that would recur under a fresh guide and a fresh donor, and does ranking targets by that dependability reorder the shortlist that effect size alone produces?
- *RQ2 (in context)*: within each activation state (Rest, Stim8hr, Stim48hr), does per-target dependability separate a context-specific signal from noise and reconstruct a coherent biological module?
- *RQ3 (the design)*: how dependable is the screen as a whole, given its two guides and four donors per target?
- *RQ4 (improving the design)*: which facet, guides or donors, limits the screen’s dependability, and how much would a future screen recover by adding to the limiting one?

RQ1 and RQ2 concern individual targets, RQ3 and RQ4 the measurement design, and each answer carries a consequence. A reordering under RQ1 would mean that effect size on its own has been nominating fragile targets and passing over reproducible ones, so the targets a screen sends for validation are the wrong ones: strong but fragile hits crowd out the reproducible ones that would actually confirm. A yes to RQ2 would be strong evidence that the coefficient measures real biology, not an artifact. If dependability reassembles a known signaling module while being told nothing about which genes belong together, the signal it scores cannot be noise, because noise does not

organize into a coherent pathway. It would also surface a kind of target that neither effect size nor a pooled ranking can see: one that is dependable in one activation state but not another. RQ3 fixes how far the present screen’s rankings can be trusted, and RQ4 turns that into guidance for the next screen, naming the facet whose reinforcement recovers the most reliability.

At the center of these questions is a per-target coefficient that ranks targets by their dependability. This is not a new concern, and dependability has been pursued along two lines. One measures reproducibility across replicates directly: the irreproducible discovery rate calls discoveries by whether they recur rather than by effect size (Li et al., 2011), and reproducibility metrics tuned to context-specific CRISPR screens show that standard quality-control does not capture whether such a hit will replicate (Billmann et al., 2023). The other decomposes expression variance into named sources, as a per-gene linear mixed model partitions variance into donor, batch, and residual components (Hoffman & Schadt, 2016); data-driven prioritization frameworks likewise integrate several weighted evidence types rather than fitness effect alone (Behan et al., 2019). Neither direction yields a per-target coefficient one can rank on. Generalizability theory joins them: it turns a per-gene variance decomposition into a single dependability coefficient for each target, on the scale of a correlation between independent repeats, and separates the facets a target must generalize across from the context it is measured within.

The rest of the paper develops the per-target coefficient and the two levels it rests on (Section 2), describes the data and the pre-registered choice of estimator (Section 3), and takes up RQ1 through RQ4 in the results (Section 4).

2 A generalizability-theory decomposition

2.1 The crossed design and its facets

We treat one perturbation’s effect as a measurement and ask how much of it would recur if the experiment were repeated. The quantity of interest is the target gene’s true effect on the transcriptome. What we observe is that effect seen through one guide, in one donor, under one activation state, and read from a finite sample of cells, and each of these can make the observation differ from the truth. This is the setting generalizability theory was built for, and the correspondence is direct (Lord & Novick, 1968): the target gene is the *object of measurement*, its true effect is the *universe score* (the value one would obtain by averaging over infinitely many independent repeats), the two guides are interchangeable versions of the same perturbation (*parallel forms*), and donors and conditions are further experimental factors, or *facets*, crossed with the target. Cell sampling and sequencing contribute the residual noise. Classical test theory writes an observation as a true score plus error, $X = T + E$ (Novick, 1966), and generalizability theory (Cronbach et al., 1972) splits that error into one term per facet,

$$E = E_{\text{guide}} + E_{\text{donor}} + E_{\text{condition}} + E_{\text{interactions}} + E_{\text{residual}},$$

and estimates how large each term is.

We distinguish random from fixed facets. Guides and donors are *random* facets: one wants the effect to generalize across the guides and donors that were not sampled. Condition is a *fixed* facet: Rest, Stim8hr, and Stim48hr are three deliberately chosen activation states, not a sample from a population of conditions, so the appropriate treatment is a separate analysis within each state rather than a generalization across states. A perturbation that acts only in the activated state is context-specific, not unreliable, and folding condition into a generalization coefficient would misclassify that signal as noise.

2.2 Per-target dependability

For target t we report dependability as a *profile* of three coefficients, one per facet, kept separate and never multiplied together. Each coefficient answers a reproducibility question: with everything held fixed except one facet, how well does the target’s transcriptional signature repeat? An observation is a (target, guide, donor, condition) pseudobulk whose effect profile spans the 18,129 genes, measured relative to a matched non-targeting control. We average a group of observations into a mean profile, and each facet coefficient is the Pearson correlation, across genes, between two mean profiles that differ only in that facet:

- R_t^{guide} : the target’s guides are split into two halves, each half’s mean profile is formed, and the two are correlated. A high value indicates that the transcriptional signature reproduces across independent guide sets.¹
- R_t^{donor} : the same construction with a two-versus-two donor split, then shrunk toward the pooled mean by a moment-based empirical-Bayes step (a global shrinkage factor in the released code, in place of a target-specific one; see Appendix A), which stabilizes a coefficient that four donors otherwise leave with only three degrees of freedom.
- $R_t^{\text{condition}}$: the mean of the pairwise correlations among the three within-state profiles, reported as cross-context *consistency* rather than as a generalization coefficient.

To rank targets we need one number rather than a profile. The per-target dependability $R_{\text{dep},t}$ is that number, and it has a concrete meaning: the correlation one would expect between two fully independent repeats of the experiment, a second run carried out with fresh guides and fresh donors. Writing an observed profile as a true signal plus independent guide and donor errors, $\tau_t + \varepsilon_{\text{guide}} + \varepsilon_{\text{donor}}$, that correlation is the signal’s share of the total variation, $\sigma_\tau^2 / (\sigma_\tau^2 + \sigma_{\text{guide}}^2 + \sigma_{\text{donor}}^2)$. The same quantity is assembled from the two split-half coefficients. Each is itself a signal share of the same kind, $R = \sigma_\tau^2 / (\sigma_\tau^2 + \sigma_e^2)$, where σ_e^2 is the error variance of the facet it refreshes. Its error-to-signal ratio is therefore $(1 - R)/R = \sigma_e^2 / \sigma_\tau^2$, and because the guide and donor errors are independent these two ratios add:

$$R_{\text{dep},t} = \frac{1}{1 + \frac{1 - R_t^{\text{guide}}}{R_t^{\text{guide}}} + \frac{1 - R_t^{\text{donor}}}{R_t^{\text{donor}}}}, \quad S_t = M_t R_{\text{dep},t}, \quad (1)$$

where M_t is the effect magnitude and S_t the dependability-weighted score used for ranking. A single-facet coefficient refreshes one source of variation; the joint coefficient refreshes both, and is therefore the most demanding form of reproducibility we report. The equality is exact when the two errors are independent, and becomes an upper bound if a guide-by-donor interaction adds a further disagreement term.

We keep a naming discipline that separates the two levels of the theory, and the two decisions it serves. At the design level, generalizability theory distinguishes the *generalizability coefficient* $E\rho^2$, for *relative* decisions such as ranking, from the *index of dependability* Φ , for *absolute* decisions against a threshold; Φ counts the facet main effects in its error as well, so $\Phi \leq E\rho^2$. Because our task is to rank targets, $E\rho^2$ is the design-level coefficient we report. The per-target quantity $R_{\text{dep},t}$ is

¹These are raw single-observation split-half correlations. The reliability of the deployed two-guide average is the Spearman–Brown step-up $2r/(1+r)$, a separate and larger quantity; we report the single-observation form because the joint coefficient is defined at the single-guide-by-single-donor level. Both are distribution-light empirical correlations rather than Gaussian variance-component ratios.

the single-target counterpart of $E\rho^2$: the generalizability of one target’s effect over the two random facets, carrying no design-level subscript. Throughout, *dependability* is the plain-language name for the trustworthiness these coefficients estimate, after Cronbach et al. (1972), whose theory of generalizability is a treatment of the dependability of measurements; the coefficients themselves keep their standard names.

3 Data and estimation

The data come from a published genome-scale Perturb-seq screen in roughly 22 million primary human CD4+ T cells (Zhu et al., 2025): CRISPR interference (CRISPRi) silences a single target gene in each cell, and single-cell RNA sequencing reads out that cell’s whole transcriptome. We do not process single cells. We start from the screen’s released *pseudobulk*, the table obtained by averaging all cells that share an experimental combination into one expression profile. Each of its 278,684 rows is one such combination, a silenced target gene, the guide RNA used to silence it, the donor the cells came from, and the T cells’ activation state (Rest, Stim8hr, or Stim48hr). It gives the average expression of 18,129 genes. The screen silences 12,731 target genes with about two guides each, crossed with four donors and those three activation states over two sequencing runs. Only about 45% of that full grid was run, and guides that target no gene serve as the controls. A row averages a median of 49 cells, and as few as one, so rows differ in how precisely they are measured, which is what sets each row’s measurement-error variance. For every row we compute a per-gene effect: the shift in each gene’s expression, in log counts per million, relative to the control cells in the same donor, condition, and run, so every effect is read against controls that share its technical context.

For each gene we ask how much of its perturbation effect is contributed by the guide facet and how much by the donor facet, the two facets a target must generalize across. A crossed mixed model, fit per gene, estimates these variance contributions and weights each pseudobulk by its precision, so rows built from few cells count for less; this is the generalizability-theory design of Section 2 (Brennan, 2001). We did not take this estimator for granted. On simulated data whose true variance components are known, we compared three candidates in a pre-registered test of which could recover the planted values, and only the precision-weighted model recovered them, under both Gaussian and heavy-tailed noise. A distance-based candidate failed: a share of a pairwise distance is not a variance, and the coefficients are defined from variances. The full estimator decision path (the candidate models and what each can and cannot identify, the two levels of the design, the per-target formulas, and a crosswalk to what the released code computes) is given in Appendix A.

4 Results

4.1 Measurement error masks dependable signal

The residual of a pseudobulk decomposition combines biological idiosyncrasy with measurement error, and the latter is inflated when a pseudobulk aggregates few cells. The non-targeting controls carry no perturbation, so the variance of the non-targeting effect within its stratum estimates the per-gene measurement-error floor directly. Subtracting this floor from the residual raises the median target-signal share from .019 to .081 and raises the number of genes with a target-signal share above .10 from 40 to 7,674, or 42% of the genome (Figure 1). The dependable signal was present but masked; the same controls give a clean non-targeting-versus-non-targeting negative test ($r \approx 0$), so the correction introduces no signal where none exists.

Removing measurement error via the non-targeting controls




Figure 1: Measurement-error removal via the non-targeting controls. Left, the median target-signal share; right, the number of genes with a dependable target-signal share above .10, on a logarithmic axis.

4.2 A reliability-weighted ranking

Ranking targets by S_t reorders the candidate list relative to effect size alone (Figure 2). Of the 100 highest-effect targets, 49 fall out of the 100 most dependable, and reproducible targets that effect size underranks rise (Table 1). These are the targets a dependability-weighted screen would nominate for validation but a strength-ranked screen would overlook.

4.3 Context-specific dependability recovers the TCR module

Analyzed within activation states, dependability separates targets that are reliable across all states from targets that are reliable in one state and not another (Figure 3). Because a within-state per-target dependability rests on roughly a third of the cells and runs lower than the pooled value, states are compared to one another rather than to the pooled number. One hundred targets are context-specific in this sense. Among them, the T-cell-receptor (TCR) signaling module, the CD3 complex (CD3D, CD3G, CD247) with the proximal kinase ZAP70 and the scaffold LAT, is reliably measurable only in activated cells, with resting dependability at the noise floor (resting $R_{dep} \approx .02-.05$; stimulated $\approx .20-.28$). The knockdown effect for CD3D grows with activation (effect magnitude 0.26, 0.57, and 1.12 across Rest, Stim8hr, and Stim48hr) and becomes dependable only upon activation. The framework reconstructs a coherent signaling module from dependability alone, without gene annotation, which is the strongest available evidence that it measures a real property of the perturbation rather than an artifact of the summary.




Figure 2: Rank by effect size alone (horizontal) against rank by the dependability-weighted score (vertical). Points below the diagonal in the lower-left quadrant are demoted (strong but unreliable); points in the upper right are promoted (reproducible).

4.4 A guide-limited design

The ranking and its context-specific structure describe the screen as it stands. How could a future screen be made more dependable? A design study projects the generalizability coefficient onto alternative numbers of guides and donors (Figure 4). On the representative signal genes the current design (two guides, four donors) reaches $E\rho^2 \approx .44$. The guide error term exceeds the donor error term by a factor of about 2.4; adding one guide raises the coefficient about three times as much as adding one donor. With two guides the coefficient is capped near .53, and no number of donors reaches .70, whereas .70 is attainable with about 15 guides. Guide-to-guide variability, rather than donor heterogeneity, is the binding constraint, which is actionable guidance for the design of subsequent screens. Repeating the design study within each activation state returns the same verdict in all three states (Table 2), so the recommendation is not an artifact of pooling; within a state the guide term dominates the donor term more sharply still, because most donor disagreement is a cross-state phenomenon.

Table 1: Targets that effect size underranks but dependability promotes. Rank is out of the 7,169 scored targets.

Target	Rank by effect	Rank by dependability	$R_{\text{dep},t}$
RPS3	79	4	.34
QARS1	53	5	.32
RPP21	124	8	.34
IMP4	96	13	.31
NMD3	158	17	.33

Table 2: Per-condition design study on the representative signal genes. The guide-limited verdict holds within every activation state.

Activation state	$E\rho^2$	Guide error	Donor error
Rest	.502	0.026	0.005
Stim8hr	.507	0.023	0.004
Stim48hr	.491	0.019	0.004
Pooled	.438	0.040	0.017

5 Reproducibility

The pipeline runs end to end from the released summary statistics, without processing the full single-cell data: data retrieval, a checksummed evidence manifest, identifiability gates, synthetic method selection, per-gene decomposition, measurement-error removal, the reliability-weighted ranking, the per-context analysis, and the design study. All results are committed as tables and figures, and every reported number regenerates from a single re-run. The code is released under Apache-2.0.

6 Limitations

Four donors leave the donor variance component with three degrees of freedom, so the donor coefficient and the donor axis of the design study carry real uncertainty, which we report as bands rather than point values. Sequencing run is partially confounded with donor, a fact of the released design that we record rather than correct. A small number of full-scale Monte Carlo gate confirmations are scheduled on a compute cluster; the results shown are the local synthetic and real-data tier. The perturbation effect is a plug-in log-CPM contrast; because the pseudobulk is raw counts, we also checked a count-native alternative that scores each effect as a multiplicative fold, a negative-binomial or Poisson log-fold-change with a library-size offset. It recovers a known effect more accurately at low counts on synthetic data, but once low-count genes are handled it leaves the target ranking essentially unchanged, because the plug-in maps a near-zero count to a near-zero effect and so already suppresses the low-count noise that would otherwise inflate sparse targets; the ranking is robust to this choice of effect estimator. Finally, this is a re-analysis of released data, so the design-study guidance is prospective, addressed to subsequent screens rather than to the present one.




Figure 3: Per-state dependability for the T-cell-receptor signaling module. Resting values sit at the noise floor; the effect becomes reliably measurable only upon activation.

A Estimator decision path

This appendix records the full path from the candidate analysis-of-variance decomposition to the per-target dependability used for ranking, so the estimator can be audited rather than taken on faith. The main text gives the readable version; here we give the design, the identifiability arithmetic, the methods we set aside, the exact formulas, and an honest crosswalk to what the released code actually computes.

A.1 Estimand, observation unit, and facets

The object of measurement is the target gene t . Its response is the effect profile $\mathbf{y}_{tgd} \in \mathbb{R}^p$ over $p = 18,129$ genes, the non-targeting-control-relative log-CPM within the matched donor-condition-run stratum. The facets are the two guides per target (nested within target; a random facet we want to generalize across), the four donors (a crossed random facet, with inference confined to the present panel), and the three activation states (a fixed facet whose levels were deliberately chosen, analyzed within level). Sequencing run is partially confounded with donor; the pseudobulk cell count sets the measurement-error variance of each observation, so profiles built from fewer cells are noisier.

A.2 The crossed model, and why it is not per-target identifiable

Written per outcome gene j , the crossed generalizability model is

$$y_{tgd,j} = \mu_j + T_{t,j} + (TG)_{tg,j} + (TD)_{td,j} + (TC)_{tc,j} + (TGD)_{tgd,j} + \varepsilon_{tgd,j},$$




Figure 4: The generalizability coefficient $E\rho^2$ across the design grid. The current screen (two guides, four donors) is marked; reaching .70 requires about 15 guides and is unreachable by adding donors.

with T, TG, TD, TGD random and the condition term TC treated as fixed (within-state). The variance components $\sigma_T^2, \sigma_{TG}^2, \sigma_{TD}^2, \sigma_{\text{res}}^2$ are *design-level*: estimable only by pooling across targets. For a single target in a single condition the design is $2 \times 4 = 8$ observations, 7 degrees of freedom: guide contributes 1, donor 3, the guide-by-donor interaction 3, leaving 0 residual. The full $2 \times 4 \times 3 = 24$ factorial is likewise saturated at its highest interaction. Without an identical-specification replicate (a second pseudobulk at the same guide, donor, condition, and run), the top-order interaction and the sampling residual cannot be separated; dropping the interaction to recover a residual is an *additive restriction* the data do not prove, not a demonstration that the interaction is zero. The $\approx 18,000$ outcome genes are correlated and do not act as independent replicates. A per-target factorial ANOVA is therefore not identifiable, which forces the two-level treatment below.

A.3 Two levels, kept distinct

Design level. Across targets, a precision-weighted crossed mixed model fit per gene by REML,

$$y \sim \text{condition} + (1 \mid \text{target}) + (1 \mid \text{target:guide}) + (1 \mid \text{target:donor}),$$

yields $\sigma_T^2, \sigma_{TG}^2, \sigma_{TD}^2$ and a residual (the reduced model absorbs the three-way interaction into the residual, per the restriction above). This is an ANOVA / mixed-model variance decomposition. The design study projects it onto alternative designs,

$$E\rho^2(n_g, n_d) = \frac{\sigma_T^2}{\sigma_T^2 + \sigma_{TG}^2/n_g + \sigma_{TD}^2/n_d + \sigma_{\text{res}}^2/(n_g n_d)},$$

the design-level coefficient of Section 4. Its residual carries measurement error; subtracting the floor estimated from the non-targeting controls is the measurement-error removal of Section 4.1.

Per-target level. A single target has only two guides and four donors, so its own variance components are unstable to non-identifiable. We therefore estimate a per-target dependability from

split-half profile reproducibility with partial pooling, not from per-target variance components. The two coefficients are not the same quantity and do not share a name: design-level is $E\rho^2$ (and Φ), per-target is $R_{\text{dep},t}$.

A.4 Candidate estimators, and why the others were set aside

The per-target method was chosen by pre-registered synthetic recovery, never by how the real ranking looked. Five candidates were weighed:

1. method-of-moments / factorial ANOVA per target: not identifiable at 2×4 (above);
2. precision-weighted crossed REML/GLS: identifiable and used at the *design* level; at the per-target level the same non-identifiability applies;
3. kernel / distance-based decomposition: a distance share is not a variance component without a positive-semidefinite-plus-additivity proof, so it is kept as a diagnostic, not a reported component;
4. robust hierarchical functional model: heavier assumptions and compute, kept as a comparator;
5. per-target split-half correlation with empirical-Bayes shrinkage: distribution-light, identifiable at 2×4 , and the only candidate to recover the planted components under both Gaussian and heavy-tailed synthetic regimes; selected.

A.5 The per-target formulas

Within condition c , let $\bar{\mathbf{y}}_{t \cdot c}$ be the target’s mean effect profile. The effect magnitude is the root-mean-square across genes,

$$M_{t,c} = \sqrt{\frac{1}{p} \sum_j \bar{y}_{t \cdot c, j}^2}.$$

The guide coefficient splits the target’s guides into halves, averages each half’s profile, and correlates them across genes,

$$R_{t,c}^{\text{guide}} = \text{cor}(\bar{\mathbf{y}}_{t, g1, \cdot, c}, \bar{\mathbf{y}}_{t, g2, \cdot, c}).$$

The correlation is Pearson, so it measures *consistency* of the transcriptional pattern, not exact agreement of magnitude; a concordance correlation would also penalize scale and location shifts, and we do not currently apply one. The donor coefficient is the same construction on a two-versus-two donor split, $R_{t,c}^{\text{donor, raw}}$, then shrunk toward the pooled mean μ_c ,

$$\tilde{R}_{t,c}^{\text{donor}} = \mu_c + w_{t,c}(R_{t,c}^{\text{donor, raw}} - \mu_c).$$

The paper-grade weight is the Fay–Herriot form $w_{t,c} = \tau_c^2 / (\tau_c^2 + s_{t,c}^2)$ with a target-specific sampling variance $s_{t,c}^2$. The released code uses a simplified global weight: it sets the within-target variance to half the total variance of the raw donor coefficients, which collapses $w_{t,c}$ to a single constant $w = 0.5$ applied to every target in the view. The target-varying Fay–Herriot weight is a planned upgrade; the constant is a documented heuristic, not the ideal estimator, and we label it as such.

Writing each coefficient as a signal share $R = \sigma_\tau^2 / (\sigma_\tau^2 + \sigma_e^2)$, its error-to-signal ratio is $(1 - R)/R = \sigma_e^2 / \sigma_\tau^2$; because the guide and donor errors are taken to be independent, the ratios add, and

$$R_{\text{dep},t,c} = \frac{1}{1 + \frac{1 - R_{t,c}^{\text{guide}}}{R_{t,c}^{\text{guide}}} + \frac{1 - \tilde{R}_{t,c}^{\text{donor}}}{\tilde{R}_{t,c}^{\text{donor}}}}, \quad S_{t,c} = M_{t,c} R_{\text{dep},t,c}.$$

Table 3: What each quantity is in the released code, against its paper-grade target.

Symbol	Released implementation	Paper-grade target	Status
$M_{t,c}$	RMS of the target mean profile	same	exact
R^{guide}	Pearson split-half across genes	(add agreement variant)	consistency only
$R^{\text{donor,raw}}$	two-vs-two Pearson split-half	same	exact
$w_{t,c}$	global 0.5 (half-variance split)	Fay–Herriot $\tau^2/(\tau^2 + s_t^2)$	heuristic; planned
R_{dep}	error-adding joint, fail-closed NA	+ interaction term	exact under additivity
$S_{t,c}$	$M_{t,c} \times R_{\text{dep}}$	same	exact
$\sigma_{T,TG,TD,\text{res}}^2$	REML <code>lmer</code> per gene	same	exact (design level)
$E\rho^2$	D-study projection	same	exact
me-removal	NTC floor from residual	same	exact (design path)
R_{2g}	reported separately	same	exact

A guide-by-donor interaction would add a further disagreement term, so this form is an upper bound when the interaction is nonzero. If either coefficient is non-positive or undefined the target is reported `not_identifiable` (a fail-closed NA), never zero-filled. The score $S_{t,c}$ is a ranking criterion, not the probability that a target is real. The deployed screen averages two guides, whose reliability is the Spearman–Brown step-up $R_{2g} = 2r/(1 + r)$; because the joint coefficient is defined at the single-guide-by-single-donor level, this step-up is reported separately and is not substituted into R_{dep} .

A.6 Pipeline and formula-to-code crosswalk

The analysis runs: joint pseudobulk $\rightarrow$ matched-NTC effects and covariance $\rightarrow$ design/rank gate $\rightarrow$ estimator comparison and synthetic recovery $\rightarrow$ guide/donor split-half profiles $\rightarrow$ donor shrinkage $\rightarrow R_{\text{dep}} \rightarrow$ effect-weighted ranking $\rightarrow$ within-condition repetition $\rightarrow$ D-study projection. The design-level path (variance components, D-study, measurement-error removal) and the per-target path (split-half, shrinkage, R_{dep}) diverge after the effects are formed; measurement-error removal acts on the design-level residual, not on the split-half coefficients, which absorb measurement error by construction.

A.7 Assumptions, costs, and falsifiers

The reported quantities rest on stated assumptions, each of which can fail: guide and donor errors are independent, so a guide-by-donor interaction would make R_{dep} an upper bound rather than an equality; the four donors are a panel, not a population, so donor generalization is to these donors; the $\approx 18,000$ genes are correlated, so the effective dimension is far below 18,000 and the split-half coefficients are noisier than an independent-gene count would suggest; Pearson measures consistency, not magnitude agreement; the non-targeting-control covariance is shared and heteroskedastic across strata; sequencing run is partially confounded with donor; and without an identical-specification replicate the top-order interaction and the residual are not separable per target. Restoring full per-target variance-component estimation would require more guides, more donors, and lane or library replicates.

An auditable, cross-model-hardened process




Figure 5: The recorded, cross-model-hardened process: a decision ledger enables adversarial review by a competing model, GPT-5.6 Sol, whose blocked verdict is resolved in the same loop. The verdict was blocked and worked through, not certified.

B A recorded, cross-model-hardened process

The methodology was developed under an auditable process (Figure 5). Each decision (the choice of estimator, the thresholds, the classification of each facet as random or fixed, and the criterion for identifiability) was recorded as a versioned issue using issue-driven development (Cheng, 2026a) and spec-driven development, both open-source tooling the author maintains. Because the reasoning was documented rather than implicit, the frozen record could be subjected to adversarial review by a competing frontier model, GPT-5.6 Sol, which returned a blocked verdict with 11 findings, three of them critical, spanning identifiability, measurement error, selection, validation leakage, and compute. Each finding was resolved within the same recorded loop: the falsification gates and controls were frozen before any real result was inspected, the estimator was selected as in Section 3, and a mis-specified reliability aggregation (a product of coefficients) was replaced by the joint coefficient of Section 2.2. The resolutions were then cross-verified across models with a second adversarial pass and with **parallel-ai-agents** (Cheng, 2026b), a package the author maintains that dispatches a task to independent agents and cross-compares their outputs. The recorded decisions, the blocked verdict, and every resolution are versioned in the repository.

The manuscript itself was produced under the same discipline. It was drafted and revised in an extended interactive session with Claude Opus, run with a one-million-token context window at maximum reasoning effort, while the author reviewed each section, formula, and claim and directed every revision. The commit history records that revision trail alongside the methodological one, so the writing, like the analysis, is auditable rather than opaque.

References

- Behan, F. M., et al. (2019). Prioritization of cancer therapeutic targets using CRISPR–Cas9 screens. *Nature*, 568, 511–516. <https://doi.org/10.1038/s41586-019-1103-9>
- Billmann, M., Ward, H. N., Aregger, M., Costanzo, M., Andrews, B. J., Boone, C., Moffat, J., & Myers, C. L. (2023). Reproducibility metrics for context-specific CRISPR screens. *Cell Systems*, 14(5), 418–422.e2. <https://doi.org/10.1016/j.cels.2023.04.003>
- Brennan, R. L. (2001). *Generalizability theory*. Springer.
- Cheng, C. (2026a). *Issue-driven development (IDD): A decision-ledger methodology for Claude Code*. <https://github.com/PsychQuant/issue-driven-development>
- Cheng, C. (2026b). *parallel-ai-agents: Parallel dispatch and cross-comparison across independent AI agents*. <https://github.com/PsychQuant/parallel-ai-agents>
- Cronbach, L. J., Gleser, G. C., Nanda, H., & Rajaratnam, N. (1972). *The dependability of behavioral measurements: Theory of generalizability for scores and profiles*. Wiley.
- Dixit, A., Parnas, O., Li, B., Chen, J., Fulco, C. P., Jerby-Arnon, L., Marjanovic, N. D., Dionne, D., Burks, T., Raychowdhury, R., Adamson, B., Norman, T. M., Lander, E. S., Weissman, J. S., Friedman, N., & Regev, A. (2016). Perturb-seq: Dissecting molecular circuits with scalable single-cell RNA profiling of pooled genetic screens. *Cell*, 167, 1853–1866.e17. <https://doi.org/10.1016/j.cell.2016.11.038>
- Hoffman, G. E., & Schadt, E. E. (2016). variancePartition: Interpreting drivers of variation in complex gene expression studies. *BMC Bioinformatics*, 17, 483. <https://doi.org/10.1186/s12859-016-1323-z>
- Li, Q., Brown, J. B., Huang, H., & Bickel, P. J. (2011). Measuring reproducibility of high-throughput experiments. *Annals of Applied Statistics*, 5(3), 1752–1779. <https://doi.org/10.1214/11-AOAS466>
- Lord, F. M., & Novick, M. R. (1968). *Statistical theories of mental test scores*. Addison-Wesley.
- Norman, T. M., Horlbeck, M. A., Replogle, J. M., Ge, A. Y., Xu, A., Jost, M., Gilbert, L. A., & Weissman, J. S. (2019). Exploring genetic interaction manifolds constructed from rich single-cell phenotypes. *Science*, 365, 786–793. <https://doi.org/10.1126/science.aax4438>
- Novick, M. R. (1966). The axioms and principal results of classical test theory. *Journal of Mathematical Psychology*, 3(1), 1–18. [https://doi.org/10.1016/0022-2496\(66\)90002-2](https://doi.org/10.1016/0022-2496(66)90002-2)
- Papalexi, E., Mimitou, E. P., Butler, A. W., Foster, S., Bracken, B., Mauck, W. M., Wessels, H.-H., Hao, Y., Yeung, B. Z., Smibert, P., & Satija, R. (2021). Characterizing the molecular regulation of inhibitory immune checkpoints with multimodal single-cell screens. *Nature Genetics*, 53, 322–331. <https://doi.org/10.1038/s41588-021-00778-2>
- Peidli, S., Green, T. D., Shen, C., Gross, T., Min, J., Garda, S., Yuan, B., Schumacher, L. J., Taylor-King, J. P., Marks, D. S., Luna, A., Blüthgen, N., & Sander, C. (2024). scPerturb: Harmonized single-cell perturbation data. *Nature Methods*, 21, 531–540. <https://doi.org/10.1038/s41592-023-02144-y>
- Replogle, J. M., Saunders, R. A., Pogson, A. N., Hussmann, J. A., Lenail, A., Guna, A., Mascibroda, L., Wagner, E. J., Adelman, K., Lithwick-Yanai, G., Iremadze, N., Oberstrass, F., Lipson, D., Bonnar, J. L., Jost, M., Norman, T. M., & Weissman, J. S. (2022). Mapping information-rich genotype-phenotype landscapes with genome-scale perturb-seq. *Cell*, 185, 2559–2575.e28. <https://doi.org/10.1016/j.cell.2022.05.013>
- Zhu, R., Dann, E., Yan, J., Reyes Retana, J., Goto, R., Guitche, R. C., Petersen, L. K., Ota, M., Pritchard, J. K., & Marson, A. (2025). Genome-scale perturb-seq in primary human CD4+ T cells maps context-specific regulators of T cell programs and human immune traits. *bioRxiv*. <https://doi.org/10.64898/2025.12.23.696273>